


Performance Report — Iris Classification

Model

Logistic regression trained on the Iris dataset (150 samples, 75/25 train/test split).

Test accuracy: 0.947 – 94.7%

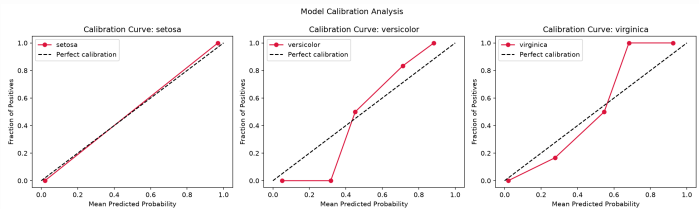
Confusion Matrix



Confusion matrix

Errors are concentrated between *versicolor* and *virginica*, the two species that are most similar in petal size.

Calibration



Calibration curves

The model is well-calibrated for *setosa*; for the other two classes, the predicted probabilities deviate somewhat more from the ideal diagonal.

Conclusion

The pipeline trains and evaluates correctly. The remaining error is explained by the natural overlap between two of the three species, also visible in the exploration notebook.